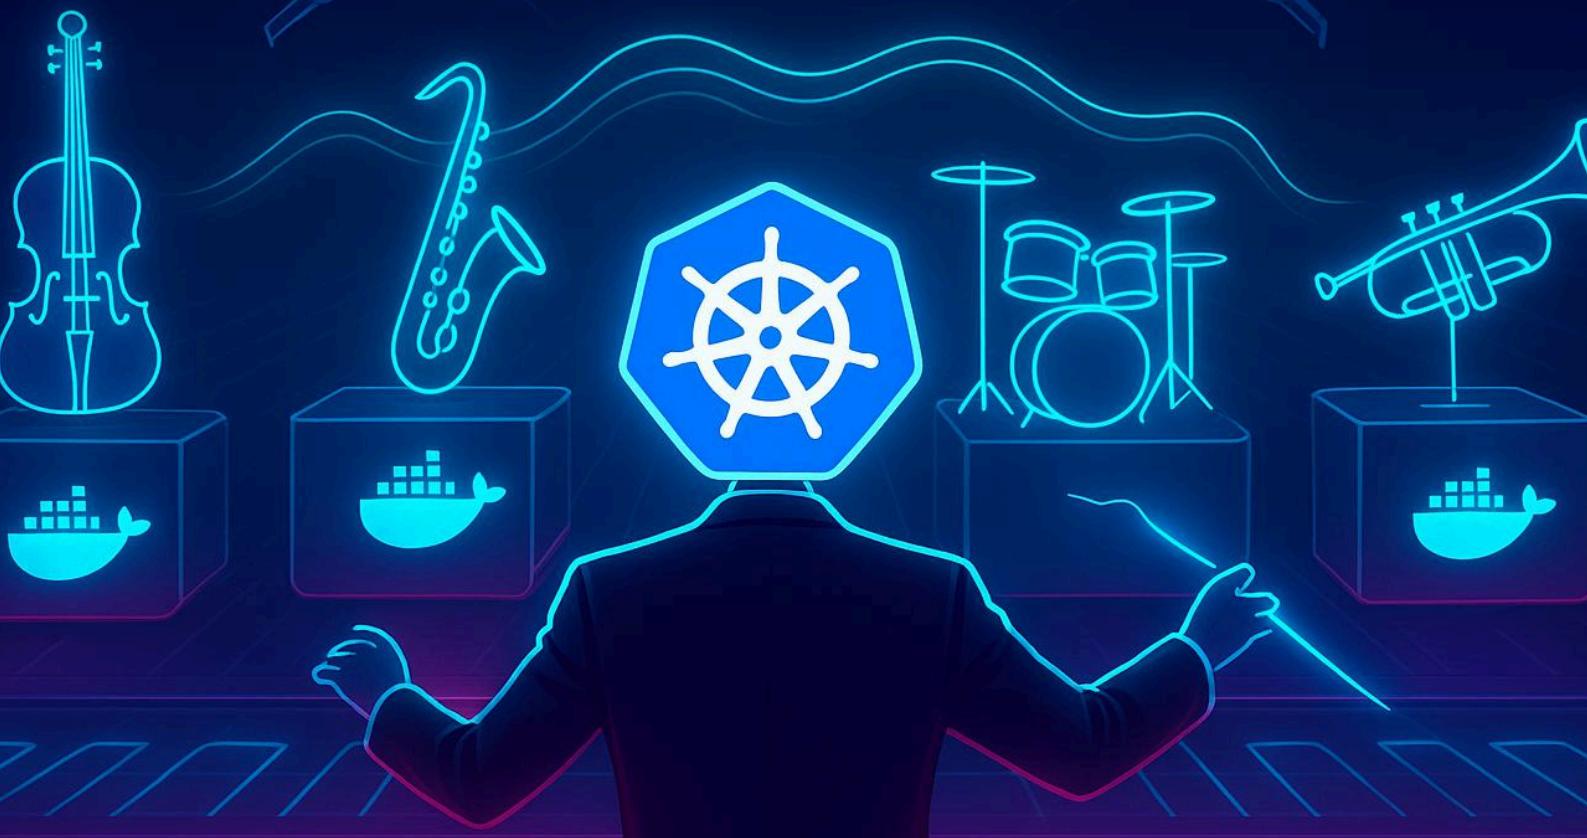




Kubernetes Handbook

Created by **JS Mastery**

Visit *jsmastery.com* for more



What's in the guide?

Kubernetes can feel overwhelming at first. There are pods, deployments, services, YAML files, and a command-line tool (`kubectl`) with what feels like a hundred flags to remember.

The goal of this guide is to make all of that a little less intimidating.

Think of this cheat sheet as your quick reference when you're working with Kubernetes day to day.

Instead of searching through docs or scrolling Stack Overflow, you'll find the most common commands, YAML patterns, and explanations right here in one place.

Whether you're just starting out or you've been running clusters for a while, this guide is built to save you time.

Bookmark it, keep it open while you work, and don't be afraid to experiment with the snippets.

...before you go

While this Kubernetes Cheat Sheet helps you quickly recall commands and YAML patterns, imagine taking the next step by learning how to design & run full production systems.

If you're eager to go beyond the basics – from building APIs and managing databases to handling authentication, scaling, and deploying on AWS – our upcoming [**Ultimate Backend Course**](#) has you covered.

The Ultimate Backend Course



It's the perfect blend of strong fundamentals and hands-on projects, guiding you from core backend principles to real-world DevOps and cloud deployment.

Join the waitlist today and be the first to get access. 

What is Kubernetes?

Kubernetes (K8s) is an open-source container orchestration platform that automates the deployment, scaling, and management of containerized applications.

Instead of manually starting and stopping containers, Kubernetes ensures that your apps are always running, healthy, and scalable.

Key Features:

- **Automatic scaling** – Scale applications up/down based on demand
- **Self-healing** – Restart failed containers and replace unhealthy nodes
- **Load balancing** – Distribute traffic across healthy containers
- **Rolling updates** – Deploy new versions with zero downtime
- **Service discovery** – Automatically expose and discover services

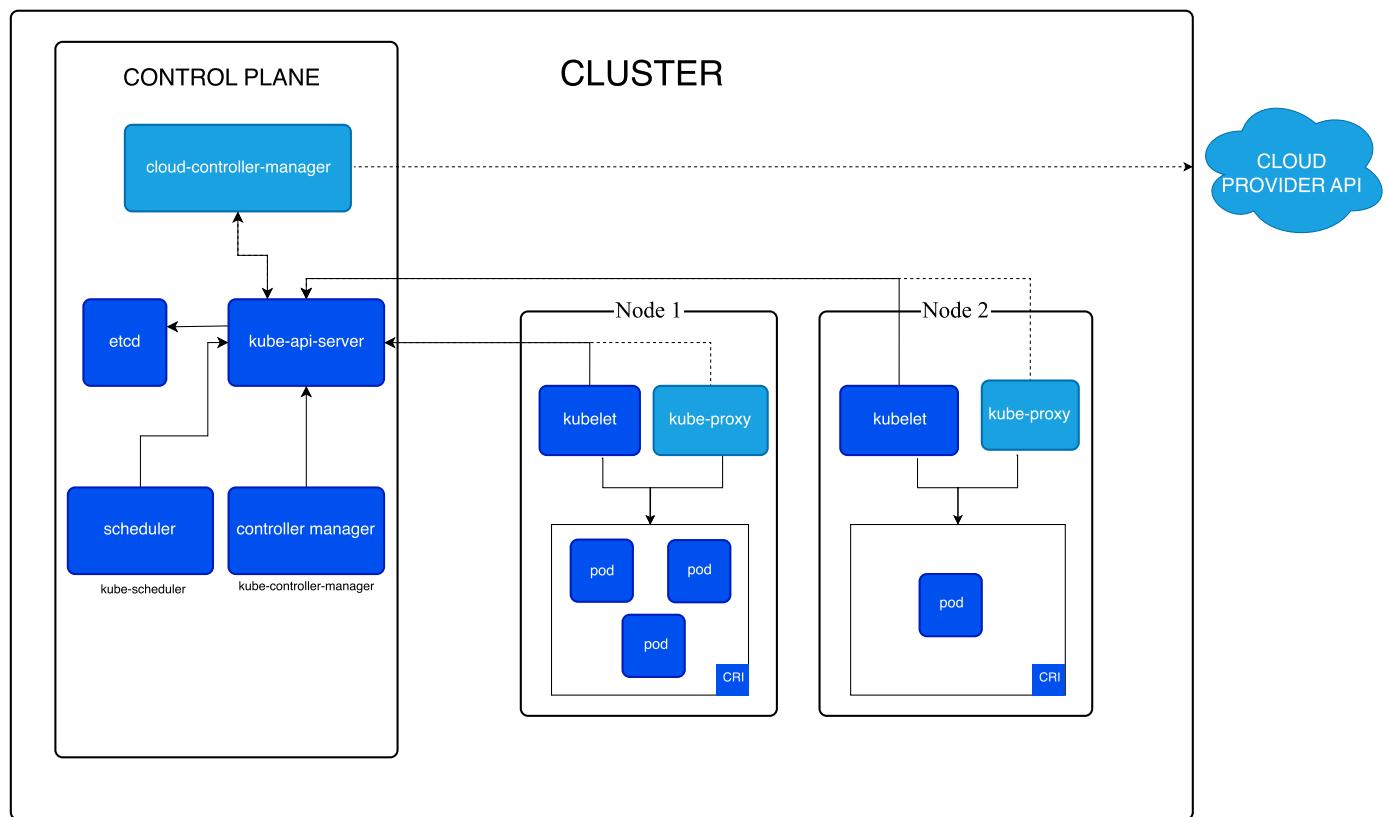
Core Components

Cluster

- A Kubernetes cluster is the full system: a set of nodes (machines) that run your workloads.

Two main parts:

- **Control Plane (brains)** → schedules workloads, maintains desired state.
- **Nodes (workers)** → actually run your apps in containers.



Cluster Components

Control Plane Components

API Server

The front door of the cluster; everything talks to it (kubectl, dashboards, other services).

etcd

A distributed key-value store that holds the cluster state (like the cluster's memory).

Controller Manager

Watches the cluster and makes changes to keep it in the desired state (e.g., if a pod crashes, it spins up a new one).

Scheduler

Decides which node a new pod should run on based on resources and constraints.

Cluster Components

Node Components

Kubelet

Agent running on each node, talks to the API server and ensures pods are running as expected.

Container Runtime

Software that runs containers (Docker, containerd, CRI-O).

Kube-proxy

Handles network rules so services can communicate with pods across nodes.

Core Components

Pod

- Smallest deployable unit in Kubernetes
- Contains one or more containers that share storage and network
- Containers in a pod share the same IP address and port space

Node

- Worker machine (physical or virtual) that runs pods
- Contains kubelet (node agent), container runtime, and kube-proxy

Deployment

- Manages replica sets and provides declarative updates for pods
- Ensures desired number of pod replicas are running
- Handles rolling updates and rollbacks

Core Components

Service

- Stable network endpoint to access pods
- Provides load balancing across pod replicas
- Types: ClusterIP, NodePort, LoadBalancer, ExternalName

ConfigMap

- Stores configuration data as key-value pairs
- Separates configuration from container images
- Can be mounted as volumes or environment variables

Secret

- Stores sensitive data (passwords, tokens, keys)
- Base64 encoded and encrypted at rest
- Similar to ConfigMap but for confidential data

Core Components

Ingress

- Manages external HTTP/HTTPS access to services
- Provides SSL termination, load balancing, and name-based virtual hosting
- Requires an Ingress Controller

Namespace

- Virtual cluster for organizing resources
- Provides scope for names and resource quotas
- Default namespaces: default, kube-system, kube-public

ReplicaSet

- Ensures a fixed number of pod replicas are running.
- Acts as the foundation for higher-level objects like Deployments.
- Automatically replaces failed or deleted pods to maintain the desired count.

Core Components

DaemonSet

Ensures one pod runs on every node (e.g., for logging/monitoring agents).

StatefulSet

Like Deployment but for apps that need stable network IDs and persistent storage (e.g., databases).

Job & CronJob

Run pods that do work and exit. CronJob schedules them like cron.

Horizontal Pod Autoscaler (HPA)

Automatically scales pods up/down based on CPU/memory usage or custom metrics.

Role-Based Access Control (RBAC)

Manages who can do what in a cluster.

YAML Manifests

Kubernetes resources are defined using YAML files with four main sections:

```
apiVersion: v1          # API version
kind: Pod               # Resource type
metadata:
  name: my-pod         # Resource metadata
  namespace: default
  labels:
    app: my-app
spec:                   # Resource specification
  containers:
    - name: my-container
      image: nginx:1.21
      ports:
        - containerPort: 80
```

Common apiVersions:

- **v1** – Core API (Pod, Service, ConfigMap)
- **apps/v1** – Deployments, ReplicaSets
- **networking.k8s.io/v1** – Ingress
- **rbac.authorization.k8s.io/v1** – RBAC

kubectl Commands

Get cluster information

```
Terminal

> kubectl cluster-info
> kubectl version
```

Get all resources

```
Terminal

> kubectl get all
> kubectl get all --all-namespaces
```

Get specific resources

```
Terminal

> kubectl get pods
> kubectl get deployments
> kubectl get services
> kubectl get nodes
```

Working with Pods

List pods

```
Terminal

> kubectl get pods
> kubectl get pods -o wide
> kubectl get pods --show-labels
```

Describe pod (detailed info)

```
Terminal

> kubectl describe pod <pod-name>
```

Get pod logs

```
Terminal

> kubectl logs <pod-name>
> kubectl logs <pod-name> -f          # Follow logs
> kubectl logs <pod-name> -c <container> # Multi-container pod
```

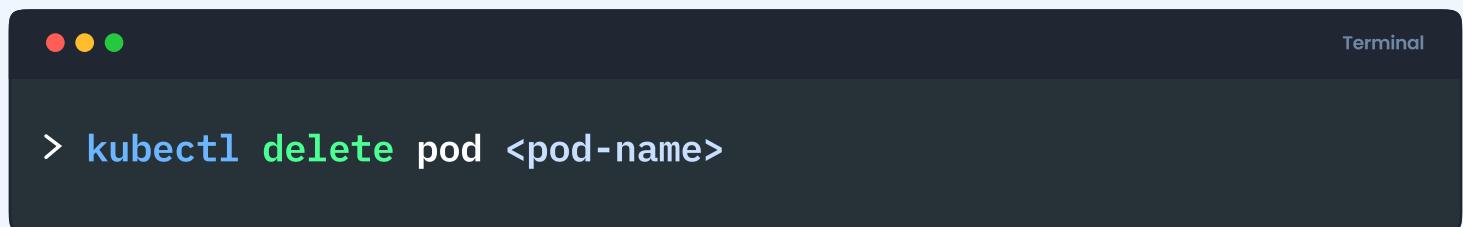
Working with Pods

Execute commands in pod



```
> kubectl exec <pod-name> -- <command>
> kubectl exec -it <pod-name> -- /bin/bash
```

Delete pod



```
> kubectl delete pod <pod-name>
```

Working w/ Deployments

List deployments

```
Terminal  
> kubectl get deployments
```

Create deployment

```
Terminal  
> kubectl create deployment <name> --image=<image>  
> kubectl create deployment nginx --image=nginx:1.21
```

Scale deployment

```
Terminal  
> kubectl scale deployment <name> --replicas=5
```

Working w/ Deployments

Update deployment image

```
Terminal

> kubectl set image deployment/<name> <container>=<image>
```

Create deployment

```
Terminal

> kubectl rollout status deployment/<name>
> kubectl rollout history deployment/<name>
> kubectl rollout history deployment/<name>
> kubectl rollout history deployment/<name>
```

Working with Services

List services

```
Terminal

> kubectl get services
> kubectl get svc      # Short form
```

Expose deployment as service

```
Terminal

> kubectl expose deployment <name> --port=80 --target-port=8080
> kubectl expose deployment nginx --port=80 --type=LoadBalancer
```

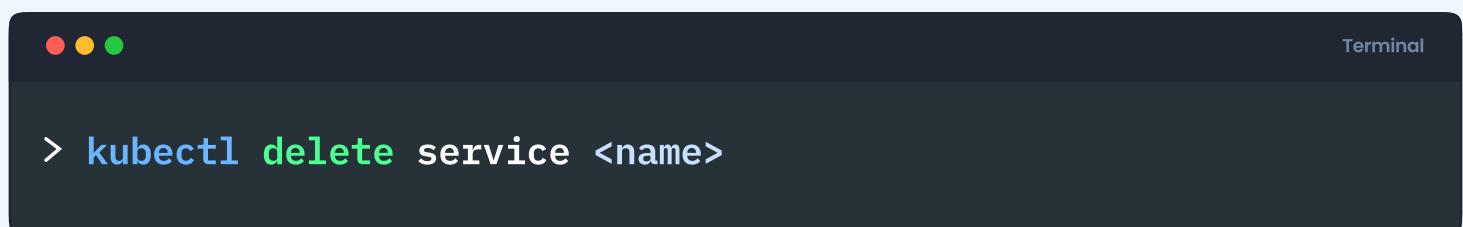
Port forwarding

```
Terminal

> kubectl port-forward service/<service-name> 8080:80
> kubectl port-forward pod/<pod-name> 8080:80
```

Working with Services

Delete service



A screenshot of a macOS terminal window titled "Terminal". The window has the standard red, yellow, and green close buttons at the top left. At the top right, it says "Terminal". The main area of the terminal shows a single line of text: "> kubectl delete service <name>".

```
> kubectl delete service <name>
```

Namespace Commands

List namespaces

```
Terminal

> kubectl get namespaces
> kubectl get ns      # Short form
```

Create namespace

```
Terminal

> kubectl create namespace <name>
```

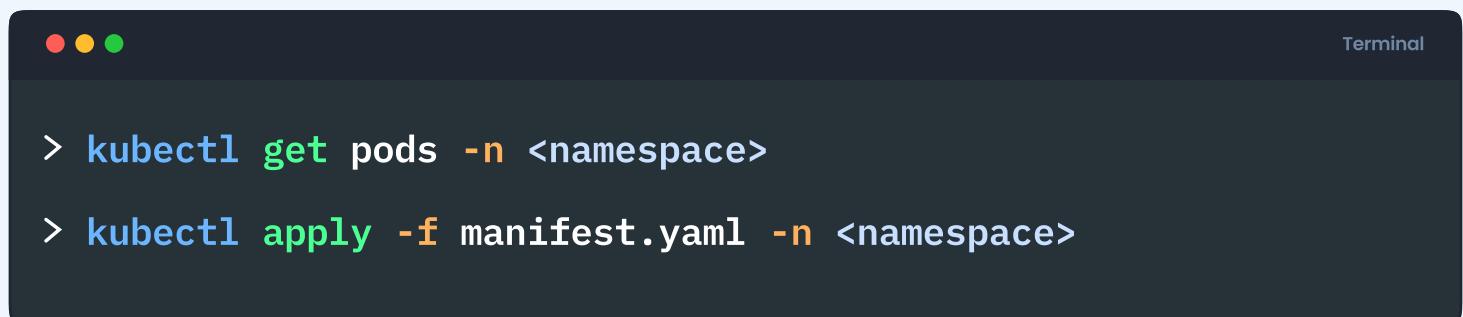
Set default namespace

```
Terminal

> kubectl config set-context --current --namespace=<name>
```

Namespace Commands

Work with specific namespace



A screenshot of a macOS Terminal window titled "Terminal". The window has three colored window control buttons (red, yellow, green) at the top left. The title bar at the top right says "Terminal". The main pane contains two lines of text, each starting with a blue right-pointing arrowhead followed by a command:

```
> kubectl get pods -n <namespace>
> kubectl apply -f manifest.yaml -n <namespace>
```

Apply & Delete Resources

Apply configuration

```
Terminal  
> kubectl apply -f <file.yaml>  
> kubectl apply -f <directory>/  
> kubectl apply -f <url>
```

Delete resources

```
Terminal  
> kubectl delete -f <file.yaml>  
> kubectl delete pod,service <name>  
> kubectl delete all --all  #Delete all resources in namespace
```

Resource Management

Creating Deployments

```
# deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.21
          ports:
            - containerPort: 80
          resources:
            requests:
              memory: "64Mi"
              cpu: "250m"
            limits:
              memory: "128Mi"
              cpu: "500m"
```

Scaling Applications

Manually scale deployment to specific replica count

```
> kubectl scale deployment nginx-deployment --replicas=5
```

Set up horizontal pod autoscaler

```
> kubectl autoscale deployment nginx-deployment --cpu-percent=50 --min=1 --max=10
```

Check horizontal pod autoscaler status

```
> kubectl get hpa
```

Rolling Updates

Update deployment with new image version

```
Terminal

> kubectl set image deployment/nginx-deployment
  nginx=nginx:1.22
```

Check the status of rollout

```
Terminal

> kubectl rollout status deployment/nginx-deployment
```

View rollout history with revisions

```
Terminal

> kubectl rollout history deployment/nginx-deployment
```

Rollbacks

Rollback to previous version

```
Terminal  
> kubectl rollout undo deployment/nginx-deployment
```

Rollback to specific revision number

```
Terminal  
> kubectl rollout undo deployment/nginx-deployment --to-revision=2
```

Perform rolling restart of all pods

```
Terminal  
> kubectl rollout restart deployment/nginx-deployment
```

Networking & Storage

SERVICE TYPES

ClusterIP (Default)

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: nginx
  ports:
  - port: 80
    targetPort: 80
  type: ClusterIP
```

Networking & Storage

NodePort

```
apiVersion: v1
kind: Service
metadata:
  name: my-nodeport-service
spec:
  selector:
    app: nginx
  ports:
  - port: 80
    targetPort: 80
    nodePort: 30080
  type: NodePort
```

Networking & Storage

LoadBalancer

```
apiVersion: v1
kind: Service
metadata:
  name: my-loadbalancer
spec:
  selector:
    app: nginx
  ports:
  - port: 80
    targetPort: 80
  type: LoadBalancer
```

Networking & Storage

Ingress

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: my-ingress
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  rules:
  - host: myapp.example.com
    http:
      paths:
      - path: /
        pathType: Prefix
      backend:
        service:
          name: my-service
        port:
          number: 80
  tls:
  - hosts:
    - myapp.example.com
    secretName: tls-secret
```

Networking & Storage

PERSISTENT VOLUMES

Persistent Volume (PV)

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: my-pv
spec:
  capacity:
    storage: 5Gi
  accessModes:
    - ReadWriteOnce
  persistentVolumeReclaimPolicy: Retain
  hostPath:
    path: /data
```

Networking & Storage

Persistent Volume Claim (PVC)

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: my-pvc
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 5Gi
```

Networking & Storage

Using PVC in Deployment

```
spec:  
  template:  
    spec:  
      containers:  
        - name: app  
          image: nginx  
      volumeMounts:  
        - name: storage  
          mountPath: /usr/share/nginx/html  
    volumes:  
      - name: storage  
    persistentVolumeClaim:  
      claimName: my-pvc
```

Troubleshooting Cmds

Check cluster status

```
Terminal

> kubectl get componentstatuses
> kubectl get nodes
```

Debug pod issues

```
Terminal

> kubectl describe pod <pod-name>
> kubectl logs <pod-name> --previous
> kubectl get events --sort-by=.metadata.creationTimestamp
```

Check resource usage

```
Terminal

> kubectl top nodes
> kubectl top pods
```

Troubleshooting Cmds

Debug network issues

```
Terminal

> kubectl get endpoints
> kubectl describe service <service-name>
> kubectl get ingress
```

Check storage

```
Terminal

> kubectl get pv
> kubectl get pvc
> kubectl describe pvc <pvc-name>
```

Useful Aliases

Add to your shell profile:

```
> alias k='kubectl'  
> alias kg='kubectl get'  
> alias kd='kubectl describe'  
> alias kdel='kubectl delete'  
> alias kaf='kubectl apply -f'  
> alias kex='kubectl exec -it'  
> alias klog='kubectl logs -f'
```

YAML Best Practices

DO:

- Use proper indentation (2 spaces)
- Always specify resource limits and requests
- Use meaningful labels and selectors
- Version your container images (avoid **latest**)
- Use ConfigMaps and Secrets for configuration

AVOID:

- Mixing tabs and spaces
- Running containers as root
- Using **latest** tag in production
- Hardcoding values in manifests
- Ignoring resource limits

Common YAML Gotchas

✗ Wrong indentation

```
spec:  
  containers:  
    - name: app
```

✓ Correct indentation

```
spec:  
  containers:  
    - name: app
```

Common YAML Gotchas

✗ Missing quotes for string numbers

```
env:  
- name: PORT  
  value: 8080
```

✓ Proper quoting

```
env:  
- name: PORT  
  value: "8080"
```

Common YAML Gotchas

✗ Incorrect selector matching

```
selector:  
  matchLabels:  
    app: nginx  
template:  
  metadata:  
    labels:  
      name: nginx    # Doesn't match!
```

✓ Matching labels

```
selector:  
  matchLabels:  
    app: nginx  
template:  
  metadata:  
    labels:  
      app: nginx
```

Quick Reference Cmds

Get everything in current namespace

```
Terminal  
> kubectl get all
```

Watch resources (real-time updates)

```
Terminal  
> kubectl get pods -w
```

Output to YAML/JSON

```
Terminal  
> kubectl top pod <name> -o yaml  
> kubectl top service <name> -o json
```

Quick Reference Cmds

Dry run (validate without creating)

```
Terminal

> kubectl apply -f manifest.yaml --dry-run=client
> kubectl create deployment test --image=nginx --dry-run=client
-o yaml
```

Force delete stuck resources

```
Terminal

> kubectl delete pod <name> --force --grace-period=0
```

Copy files to/from pods

```
Terminal

> kubectl cp <local-path> <pod>:<remote-path>
> kubectl cp <pod>:<remote-path> <local-path>
```

Resource Monitoring

Check cluster status

```
Terminal

> kubectl get componentstatuses
> kubectl get nodes -o wide
```

Debug pod issues

```
Terminal

> kubectl describe pod <pod-name>
> kubectl logs <pod-name> --previous #Previous container logs
> kubectl get events --sort-by=.metadata.creationTimestamp
```

Debug network issues

```
Terminal

> kubectl get endpoints
> kubectl describe service <service-name>
> kubectl describe ingress <ingress-name>
```

Resource Monitoring

Check storage

```
Terminal

> kubectl get pv
> kubectl get pvc
> kubectl describe pvc <pvc-name>
> kubectl get storageclass
```

RBAC debugging

```
Terminal

> kubectl auth can-i create pods
```

Maintenance Commands

Uncordon node after maintenance

```
...
```

Terminal

```
> kubectl uncordon <node-name>
```

Cordon node (prevent new pods)

```
...
```

Terminal

```
> kubectl cordon <node-name>
```

Taint node

```
...
```

Terminal

```
> kubectl taint nodes <node-name> key=value:NoSchedule
```

Remove taint

```
...
```

Terminal

```
> kubectl taint nodes <node-name> key:NoSchedule-
```

The End

🎯 Remember: Kubernetes is a Marathon, Not a Sprint

Every resilient cluster starts small. Don't try to master every resource at once — begin with pods and deployments, grow into services and ingress, and layer on scaling and security as you go.

Your future self (and your team) will thank you for investing in a solid foundation.

Happy shipping! 🚀

P.S. If this guide sparked your curiosity about running real-world systems, you'll love our upcoming [Ultimate Backend Course](#). Join the waitlist today for early access.

**The Ultimate
Backend Course**

